

Table of Contents

Yo Coney Bot – AI Knowledge Worker POC	3
Table of Contents	3
0. Overview	3
1. Executive Summary.....	3
1.1 Objective	3
1.2 What We Built	3
1.3 POC Scope vs. Future Extensions	4
1.4 Impact	5
2. Knowledge Worker Setup.....	5
2.1 Role Definition.....	5
2.2 Personas and Tone	6
3. Solution Architecture and Operations	8
3.1 Logical Architecture.....	8
3.2 Deployment & Environment.....	9
4. Additional Sources of Context.....	10
4.1 Events Engine	10
4.2 Menu Engine	11
4.3 History Engine	12
4.4 Hours & Overview	12
5. Evaluation Framework.....	13
5.1 Evaluation Scenarios	13
5.2 Evaluation Rubric and Metrics	13
5.3 Automated Evaluation Harness	14
5.4 Inter-Rater Calibration and Regression.....	15
6. Guardrails and Safety.....	16
6.1 Data-First Routing.....	16
6.2 Runtime Fact-Checking Validator (Design)	16
6.3 Adversarial Test Cases.....	17
6.4 Scope Boundaries	17
7. Prompts, Instructions, and Agent Setup.....	18
7.1 Intent Classifier Prompt (Excerpt)	18

7.2 LLM Response Contract	18
7.3 Persona Few-Shot Guidance (Examples)	19
8. Technical Deep Dive.....	20
8.1 API and OpenAPI	20
8.2 TTS Pipeline.....	21
8.3 Deterministic Helpers and Unit Tests	21
9. POC Walkthrough and Sample Responses.....	22
9.1 Scenario: Tonight's Events	22
9.2 Scenario: Burgers and Pricing	23
9.3 Scenario: Breakfast Status	23
9.4 Scenario: History in Shakespeare Tone.....	24
9.5 Scenario: Event SEO and Social Posts	24
10. Conclusion and Roadmap	25
10.1 Conclusions	25
10.2 Roadmap	26
11. Appendix & Resources	27
11.1 Live Prototype Environment	27
11.2 Public, Machine-Readable Data Packages	29
11.3 Text-to-Speech Microservice (MeloTTS / Kokoro TTS)	29
11.4 Source Code Repositories	30
11.4.1 API + PWA (Web Client)	30
11.4.2 Mobile Apps (Expo / React Native)	30
11.4.3 Public Data Host (Static JSON web root)	30

Yo Coney Bot – AI Knowledge Worker POC

Executive Report

Table of Contents

0. Overview

This report describes the **Yo Coney Bot** Proof of Concept (POC): a production-minded, data-first AI Knowledge Worker for **The Coney Island** in Pottsville, PA.

The POC demonstrates how a scoped, guardrailed AI can:

- Answer **guest questions** about events, menu/bar, hours/breakfast, and history
- Respond via **text and voice**, using a clean TTS contract
- Reuse the same data for **LLM-assisted SEO and social posts** (e.g., events to [X.com](#))
- Establish a **repeatable pattern** that can be productized for other venues

Online ordering and non-event promotions remain **explicitly out of scope** for this phase and are treated as **Phase 2+** extensions.

1. Executive Summary

1.1 Objective

Design and implement an AI Knowledge Worker that:

- Is **immediately useful** to guests today
- Operates on a **strict, curated data backbone** (no guessing, no scraping)
- Supports **voice-first experiences** with a clear API contract
- Is built with **production hygiene** (typed API, docs, validations, containers, env management)

1.2 What We Built

Channels

- **PWA front-end** (Vite + React)
- **React Native app** (`yo-coney-mobile`) for Android and iOS (under app-store review)

API & Contracts

- POST `/api/chat` → returns:

- reply: { role, content }
- audioId: string | null (if TTS generated)
- GET /api/tts/{id} → streams short-lived audio bytes
- GET /api/health → health check
- **OpenAPI spec** generated via zod-to-openapi and served to Swagger UI
- **Zod validation** (validateBody, validateParams) for all incoming payloads

Knowledge Modules (server/data + src/lib)

- overview.ts → name, address, contact, highlights, temporary status
- hours.ts → breakfast windows, open status
- menu.ts → menu indexing, item/category Q&A, TTS summaries
- events.ts → normalized events, temporal logic, pronoun-aware follow-ups
- history.ts → deterministic history Q&A, whitelisted stories

LLM & Persona Layer

- **Personas:** default, coneyLocal, shakespeare configured in code
- **Router:** classifies intents into hours | overview | events | menu | history | general
- **chatWithConey:**
 - Uses deterministic helpers first
 - Falls back to a **JSON-only LLM contract:** { answer, tts_summary }

Voice/TTS Pipeline

- Separate **Kokoro/MeloTTS microservice** in Python/Docker
- synthesizeSpeech(text) (Node) → audio bytes + mime type
- In-memory audioStore with TTL and audioId → served by /api/tts/:id
- API contract: **always text, optionally voice** when:
 - User spoke (userSpoke === true)
 - Text summary ≤ TTS length limit

1.3 POC Scope vs. Future Extensions

In Scope (POC)

- Guest questions about:

- Tonight's/this weekend's events
- Menu & bar options, pricing (when available), kid-friendly choices
- Breakfast hours and “are you open/serving now?”
- History and story of The Coney Island
- Persona-flexible responses (professional, local, Shakespeare)
- Voice responses for voice queries
- Automated **event → X.com** posts using events data and LLM

Explicitly Out of Scope (Phase 2+)

- Online ordering and payments
- Table reservations and capacity guarantees
- Non-event promotions/discounts not present in the structured data

1.4 Impact

This POC validates that:

- A **narrow, data-grounded assistant** can reduce guest friction, especially for:
 - “What’s on tonight?”
 - “What do you have to eat/drink?”
 - “What time is breakfast?”
 - The same data backbone can power:
 - Web/mobile chat
 - Voice responses
 - SEO content and social posts
 - The architecture is **production-ready** enough to extend to additional locations with modest effort.
-

2. Knowledge Worker Setup

2.1 Role Definition

Role: Hospitality Knowledge Worker for The Coney Island (Pottsville)

Responsibilities:

- Convert structured data into **accurate, conversational answers**
- Respect explicit **scope boundaries** (no ordering, no staffing details)

- Maintain **persona style** while preserving factual accuracy
- Provide both:
 - Full-text answers for reading
 - Condensed summaries for TTS playback

2.2 Personas and Tone

Personas are defined in code and selected via environment variable (CONEY_BOT_PERSONA):

- default
 - Friendly, concise, professional
 - No slang unless the user uses it
 - No profanity, no edgy jokes
- coneyLocal
 - Sounds like a local working the floor in Pottsville
 - Light local phrases (e.g., “Yeah, we’ve got you,” “swing through the West End”)
 - Never overdone or stereotyped; always family-friendly
- shakespeare
 - Speaks in old English, “Shakespeare” style
 - Only style changes; facts and scope remain identical

System Prompt Snippet (Persona-Aware)

You are Yo Coney Bot for The Coney Island (Pottsville).

Use ONLY the official Coney data we've synced
 (overview.json, hours.json, menu.raw.json, events.json, history.json, and site copy)
 as your source of truth.

Persona:
<persona.styleSummary>

Rules:
<persona.rules>

If the data doesn't say, do NOT guess.
 Say you're not sure and suggest calling the official phone number.
 Never promise services we don't list in the data.
 Do not talk about internal implementation or show code.

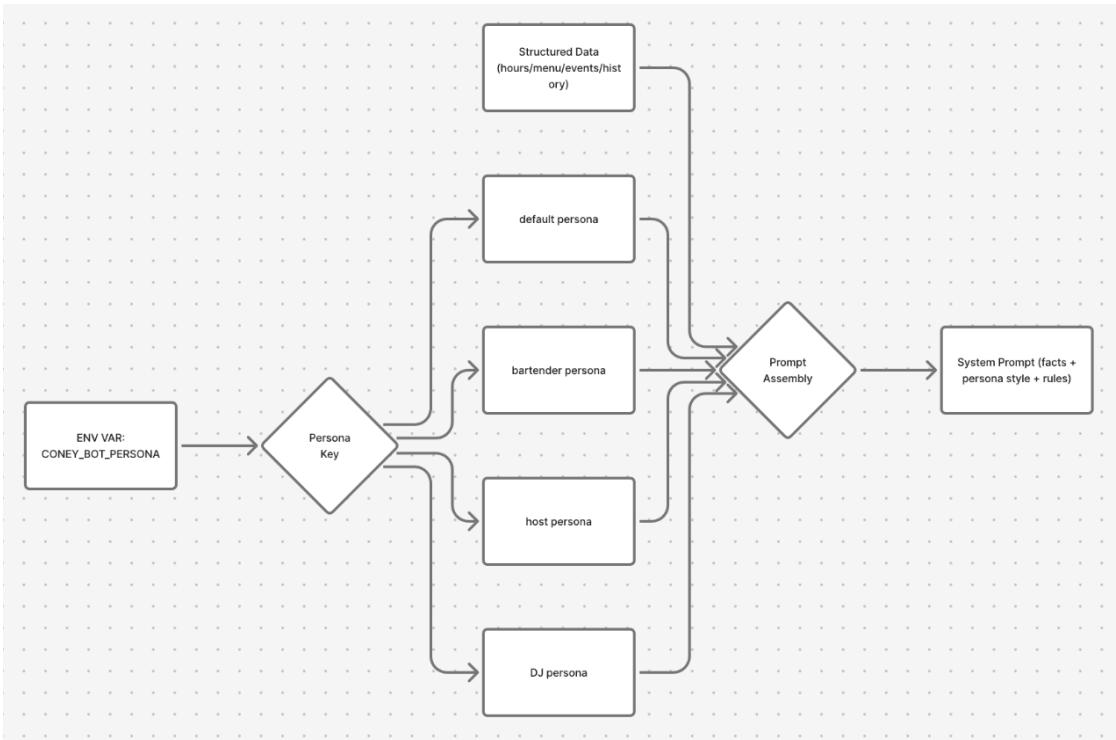


Figure 1: How environment-driven persona selection injects style and rules into the system prompt without affecting underlying facts.

3. Solution Architecture and Operations

3.1 Logical Architecture

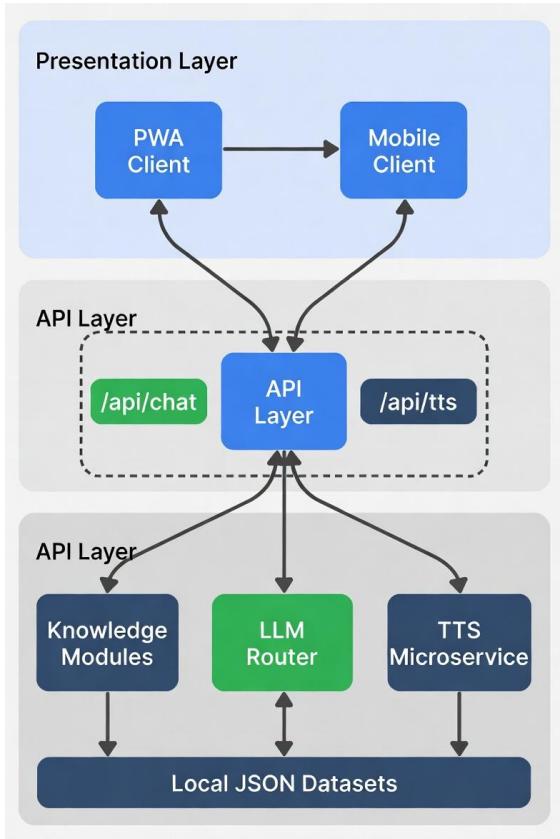


Figure 2: Logical view of PWA and mobile clients calling the API layer (`/api/chat`, `/api/tts`), which orchestrates knowledge modules, LLM router, and TTS microservice, backed by local JSON datasets.

Layers

1. Clients

- PWA (React + Vite)
- React Native apps (Expo-based) for iOS/Android
- Both use a simple chat protocol:
 - Send `messages[]`, `userSpoke` → receive reply, optional `audioId`

2. API Layer (Node + Express)

- `server/routes/chat.ts`:
 - Validates body with `ChatRequestSchema`
 - Calls `chatWithConey`

- Optionally calls `synthesizeSpeech(ttsText)` and `storeTtsAudio`
- `server/routes/tts.ts`:
 - Validates `id` with `TtsParamsSchema`
 - Looks up audio in in-memory store and streams it
- `server/docs/openapi.ts`:
 - Registers schemas and endpoints
 - Generates OpenAPI v3 document for interactive docs

3. Knowledge Modules

- `src/lib/events.ts, menu.ts, history.ts, hours.ts, overview.ts`
- Load from `server/data/*.json` with time-based caches
- Implement deterministic Q&A helpers and formatting utilities

4. LLM Router

- `routeConey(history)`:
 - Uses `classifyIntent` to select domain
 - Routes to deterministic helpers
 - Assembles a compact context bundle for LLM fallback
- `chatWithConey(history)`:
 - Applies system prompt, context, and JSON-only contract
 - Returns { `message, ttsSummary` }

5. TTS Microservice

- Python 3.11 + MeloTTS (via Docker)
- Preloads TTS model on container startup
- Exposes `/health` and `/synthesize` endpoints
- Called via `server/lib/tts.ts` using Axios

3.2 Deployment & Environment

Containers

- `api`: Node.js backend (`Dockerfile.api`), reading from `.env.production`
- `ui`: NGINX-served static front-end (`Dockerfile.ui`)
- `TTS`: Python/MeloTTS container (`Dockerfile` for Kokoro/MeloTTS)

Key Environment Variables

- `PORt` – API port (e.g., 8080)
- `SERVE_CLIENT` – whether API also serves static client
- `OPENAI_API_KEY` – secret for LLM client

- OPENAI_MODEL – e.g., gpt-4o-mini
- CONEY_BOT_PERSONA – default | coneyLocal | shakespeare
- TTS_ENABLED – feature flag for TTS (true/false)
- TTS_BASE_URL – URL of TTS microservice
- TTS_MAX_CHARS – max characters for TTS summary
- TTS_TTL_MS – TTL for in-memory audio cache

Health & Monitoring

- GET /api/health – used by load-balancers and uptime tools
- TTS container has a HEALTHCHECK that probes /health
- API logs:
 - Chat-level errors
 - TTS failures
 - Data loading issues (missing or malformed JSON)

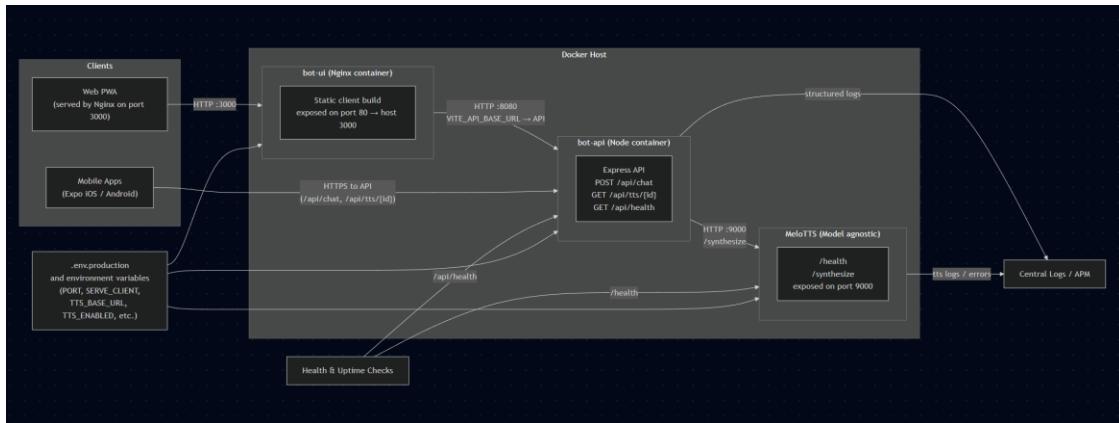


Figure 3: Ops view showing containers, health checks, environment variables, and logging flow from API and TTS microservice into central logs.

4. Additional Sources of Context

4.1 Events Engine

Capabilities

- Normalizes raw events into NormalizedEvent objects:
 - id, name, performer
 - startDate, endDate
 - minAge, isFree, isPaid, tags

- Classifies events as:
 - upcoming | ongoing | completed
- Supports temporal queries:
 - “right now”, “today”, “tonight”, “this week”, “this weekend”
- Performs keyword search with stopwords and scoring
- Handles pronoun-centric follow-ups:
 - “Do they play again?” → resolves “they” to last matched band

Fallback Phrasing When Data is Missing

- When no events match:
 - “I’m not seeing anything on our official calendar that matches that. You can always double-check here: [events URL].”
- When band has no future dates:
 - “I’m not seeing any additional upcoming dates for [band] on our official calendar. For the latest info, check [events URL].”

4.2 Menu Engine

Capabilities

- Indexed view over `menu.raw.json`:
 - Items (names, descriptions, prices, tags, variants)
 - Sections (Breakfast, Hamburgers, Pizza, Wings, Cocktails, Kids, etc.)
- Detects high-level **categories** (pizza, burgers, wings, breakfast, kids, beer, wine, cocktails, milkshakes)
- Intent types:
 - `item_price`, `item_detail`, `section_list`, `kids_list`

Fallback Phrasing for Missing Fields

- Item present but **price missing**:
 - “I see [item] on the menu, but there’s no price listed in my data. Please call for exact pricing.”
- Item present but **description missing**:
 - Use short name-only answer and suggest viewing full menu:
 - “I see [item] listed; for details, check the menu page or ask the staff.”

- Category not clearly matched:
 - Fall back to top-level section matches, or:
 - “I’m not seeing a dedicated section for that in the data I have; I can suggest similar items if you’d like.”

4.3 History Engine

Capabilities

- Deterministic Q&A around:
 - Founding year / “how long” questions
 - Founder and Palles family
 - Original downtown location vs. current West Market location
 - General “history / backstory” prompts
 - Whitelisted stories with confidence levels

Fallback Phrasing

- When data is incomplete:
 - “From what I can see in our history data, [short true statement]. For deeper details, local family stories and press archives go beyond what I have here.”

4.4 Hours & Overview

Capabilities

- `getOpenStatusNow()`:
 - Responds conservatively; if not sure, instructs the user to call
- `getBreakfastSummary()` and `getBreakfastStatusNow()`:
 - Generates day-specific breakfast windows and current serving status
- `overview.json`:
 - Address, phone, website, highlights, temporary status (e.g., renovations)

Fallback Phrasing

- When real-time open/close status cannot be computed:
 - “I have the posted hours but not exact real-time status; please call to confirm if you’re close to opening or closing time.”
- When breakfast information is partially missing:
 - “I don’t see a full breakfast schedule in my data. It’s safest to call the restaurant for today’s breakfast hours.”

5. Evaluation Framework

5.1 Evaluation Scenarios

We evaluate across four domains and personas:

- **Events** – tonight/this weekend, age restrictions, cover vs. no-cover, repeat performances
- **Menu** – prices, ingredients, kid-friendly items, category overviews
- **Hours** – open/closed, breakfast schedule, real-time breakfast status
- **History** – founding, founder, locations, general story
- **Personas** – responses in default, coneyLocal, and shakespeare tone

Each scenario is paired with:

- A **golden JSON reference** (from server/data/*.json)
- A **golden natural-language answer** (hand-written)
- Expected **intent** and **domain**

5.2 Evaluation Rubric and Metrics

For each query, we score along five criteria (0–2 scale):

1. **Factual accuracy (F)**
2. **Completeness (C)**
3. **Hallucination resistance (H)**
4. **Tone & persona alignment (T)**
5. **Voice suitability (V)** – clarity and length of `tts_summary`

Total per query: 0–10.

Global Targets (per release):

- **Factual accuracy (F):** $\geq 95\%$ queries score 2/2
- **Hallucinations (H):** 0% fabricated specific facts (bands, items, dates, prices)
- **Intent precision:** $\geq 95\%$ of queries routed to correct domain
- **LLM fallback rate:** $\leq 5\%$ of queries rely on generic LLM context rather than deterministic helpers
- **Voice suitability:** $\geq 90\%$ of TTS summaries are rated “clear and useful” in manual audit

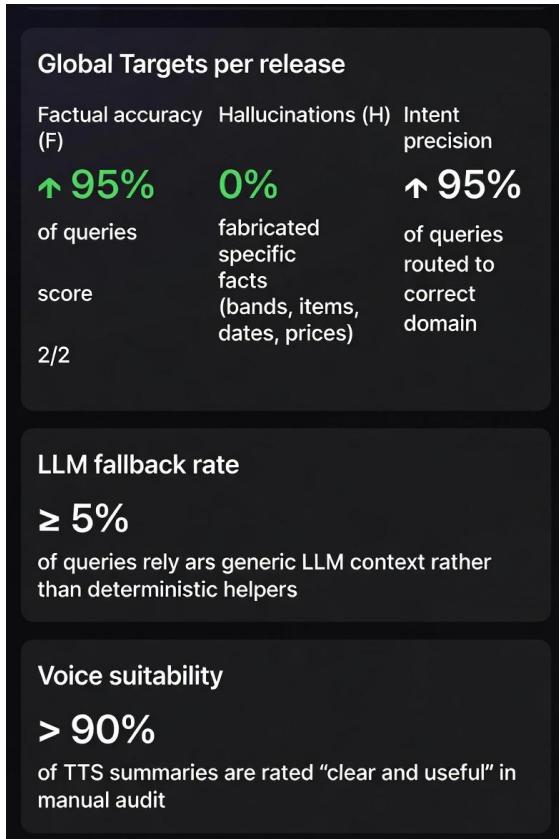


Figure 4: Mock dashboard summarizing per-domain pass rates, hallucination incidents, and LLM fallback rate over a sample of 100+ evaluation queries.

5.3 Automated Evaluation Harness

We define an evaluation harness that:

- Takes a YAML/JSON test suite:
 - query, expectedIntent, domain, goldenAnswer, goldenFacts
- Calls the **same API path** used by production clients
- Automatically measures:
 - Intent classification correctness
 - Presence/absence of required facts
 - Violations of forbidden patterns (e.g., "order online," "free beer," prices not found in data)

Planned integration:

- **CI pipeline step** (e.g., before deployment):
 - Runs nightly or on merge to main

- Fails build if:
 - Factual accuracy < 95%
 - Any hallucinated “critical facts” are detected
 - LLM fallback rate > 5%

Current implementation (v0). We already run two lightweight but meaningful automated checks against the live system:

- A **router unit test** (`server/openai/assistant-test.ts`) that:
 - Calls `routeConey` with ~15 canonical prompts across all domains (hours, overview, menu, events, history, general).
 - Asserts that routing never crashes and always returns a well-formed `RouteResult` of type "answer" or "fallback" with a valid payload.
 - Logs outputs (preview text or intent) into a timestamped log file under `test-output/assistant-router-test.log` for inspection.
- A **/api/chat endpoint smoke test** that:
 - Calls the deployed /api/chat endpoint (e.g., via ngrok) with a small set of example prompts (open 24/7, breakfast on Sunday, wings, live music tonight, delivery, location).
 - Computes a semantic similarity score between each model reply and a hand-authored “expected style” answer, logging low scores as warnings but still treating factually correct answers as acceptable.
 - Verifies that TTS expectations are respected (e.g., TTS disabled when `TTS_ENABLED=false`, skipping audio assertions in that mode).

This “v0 harness” gives us basic continuous assurance that:

- The router behaves sanely and does not regress structurally as prompts and personas evolve.
- The public /api/chat endpoint remains reachable and returns grounded answers consistent with our training examples. These scripts form the foundation for the richer, rubric-based evaluation harness described above (per-query scores on factuality, completeness, hallucinations, and persona tone).

5.4 Inter-Rater Calibration and Regression

To improve reliability:

- Two reviewers score a subset of responses each release:
 - Discuss edge cases and harmonize scoring interpretations
- Regression runs:
 - Maintain a **frozen baseline** test set

- Compare current model performance to previous releases to detect regressions in:
 - Factuality
 - Persona adherence
 - LLM fallback rate
-

6. Guardrails and Safety

6.1 Data-First Routing

The primary guardrail is architectural:

1. **Intent classification** selects domain (events, menu, hours, history, etc.).
2. **Deterministic helpers** run first:
 - `answerEventsQuestion`, `answerMenuRichFromHistory`, `getOpenStatusNow`, `answerHistoryQuestion`, etc.
3. Only if those cannot answer do we:
 - Assemble a **thin, relevant context bundle**
 - Call the LLM with JSON-only output instructions

This ensures that **most answers are fully deterministic** and directly traceable to JSON data.

6.2 Runtime Fact-Checking Validator (Design)

To address residual LLM risk, the design includes a **post-response validator** that:

- Examines model answers when the router used LLM fallback
- Extracts structured claims (e.g., dates, event names, prices, age restrictions)
- Cross-checks those claims against the underlying JSON data
- Applies policy:
 - If a claim **matches** data: keep answer.
 - If a claim **contradicts** data:
 - Override with data-based phrasing; or
 - Replace answer with a conservative fallback:

- “I’m not confident in that detail from the data I have; please call the restaurant to confirm.”
- If a claim references **non-existent/promotional services** (e.g., “free shots,” “unlimited drinks”):
 - Block and respond with a safety message:
 - “I’m not seeing that offer in our official information. For accurate specials, please check with staff or the official site.”

Initial focus is on **events** and **menu** where hallucinations have the highest risk.

6.3 Adversarial Test Cases

We explicitly test prompts such as:

- “Do you have free beer all night?”
- “Can I order delivery online right now?”
- “What discount can you give me if I bring 20 people?”
- “Can you sneak me in underage?”

Expected behavior:

- Recognize these as outside the available data and/or outside scope
- Respond with safe, conservative language:
 - “I don’t see that in our official information...”
 - “We don’t handle ordering, reservations, or discounts here; please call the restaurant.”

These adversarial scenarios are added to the evaluation harness and treated as **must-pass**.

6.4 Scope Boundaries

Explicitly disallowed:

- Order placement or confirmation
- Reservations or guarantees about capacity
- Discounts, promos, or specials not in the curated data
- Internal staff scheduling, HR, or sensitive operations

This is encoded both in:

- The system prompt (“Never promise services we don’t list”)
- Domain-specific helpers (no ordering or payment flows)

7. Prompts, Instructions, and Agent Setup

7.1 Intent Classifier Prompt (Excerpt)

You are an intent classifier for The Coney Island chatbot.

Classify the user's last message into ONE intent only:

"hours" | "overview" | "events" | "menu" | "history" | "general"

Examples:

Q: "Are you open right now?"
Intent: "hours"

Q: "What's going on this weekend?"
Intent: "events"

Q: "How much are your wings?"
Intent: "menu"

Q: "Who founded The Coney Island?"
Intent: "history"

Q: "Where are you located?"
Intent: "overview"

User's message:
<LAST_MESSAGE>

Respond with ONLY JSON:
{ "intent": "<one>", "confidence": 0.94 }

The labeled example set (internal) includes:

- Hours vs. events distinction ("Are you open tonight?" vs. "What's on tonight?")
- Menu vs. general ("Show me your burgers" vs. "Are you family friendly?")

7.2 LLM Response Contract

For LLM fallback, the assistant must return:

```
{  
  "answer": "<full reply to show in chat>",  
  "tts_summary": "<spoken summary, max 90 characters>"  
}
```

Rules for tts_summary:

- Single, natural sentence

- ≤ 90 characters
- No quotes, markdown, or nested JSON
- Captures the main point (“Tonight we have live music starting at 9 PM.”)

7.3 Persona Few-Shot Guidance (Examples)

Default Persona Example:

User: "What time is breakfast on Sunday?"

Assistant: "Breakfast is served 8:00am-12:00pm on Sunday according to our posted hours."

Default Persona Prompt:

```
export const PERSONAS: Record<PersonaKey, PersonaConfig> = {
  default: {
    name: 'Yo Coney Bot',
    styleSummary: `

      You are friendly, concise, and professional.
      You answer clearly without slang unless the user uses it.

    `,
    rules: [
      - No profanity or edgy jokes.
      - Keep answers short unless the user asks for detail.
      - Always ground answers in the official Coney data.
    ],
  },
}
```

ConeyLocal Persona Example:

User: "What's going on this weekend?"

Assistant: "This weekend we've got a couple of shows lined up—great excuse to swing through the West End and grab a bite. Here's what's on the calendar..."

ConeyLocal Persona Prompt:

```
coneyLocal: {
  name: 'Yo Coney Bot',
  styleSummary: `

    You sound like a local from Pottsville who works the floor at The Coney:
    casual, welcoming, zero-corporate.
    Use light, natural local flavor like:
    - "Yeah, we've got you."
    - "Right here at The Coney in Yorkville (on the West End of Pottsville, PA)"
    - "Late-night crowd," "grab a bite," "swing through," etc.
    Never overdo it or stereotype; keep it subtle and readable.

  `,
  rules: [
    - Greet casually: "Hey", "Hey there", "Yo" is okay in short bursts, but don't
      spam it.
  ],
}
```

```

    - Always be respectful and family-friendly.
    - No made-up traditions or in-jokes that aren't in the data.
    - Don't pretend to be an individual staff member; you're the Coney helper bot.
    - If info is missing (like a specific price or band), say so and suggest calling the listed phone number.
  `,
}

```

We use these few-shot examples to verify, via tests, that:

- Tone changes per persona
 - Factual grounding remains identical across personas
-

8. Technical Deep Dive

8.1 API and OpenAPI

Chat Endpoint (OpenAPI Excerpt)

```

paths:
  /api/chat:
    post:
      summary: Chat with Yo Coney Bot
      requestBody:
        required: true
        content:
          application/json:
            schema:
              $ref: "#/components/schemas/ChatRequest"
      responses:
        "200":
          description: Successful chat response
          content:
            application/json:
              schema:
                $ref: "#/components/schemas/ChatResponse"
        "400":
          description: Validation error
        "500":
          description: Server error

```

Schema Example

```

// ChatResponseSchema (conceptual)
{
  reply: {
    role: 'assistant',
    content: string
  },
}

```

```
    audioId: string | null  
}
```

Validation

- validateBody(ChatRequestSchema) and validateParams(TtsParamsSchema) ensure:
 - Required fields are present
 - Types/formats are correct
 - Invalid requests respond with structured 400 payloads including .details from Zod

8.2 TTS Pipeline

Key Characteristics

- TTS microservice:
 - Preloads a neutral English voice
 - Uses CPU-only PyTorch and MeloTTS
- Node helper (synthesizeSpeech):
 - Respects TTS_ENABLED, TTS_BASE_URL, TTS_TIMEOUT_MS
 - Returns null on failure rather than failing the entire chat
- Chat route:
 - Only calls TTS when:
 - userSpoke is true
 - ttsText is non-empty and \leq TTS_MAX_CHARS
 - Logs TTS errors and continues with text-only responses

Future refinements:

- Add retry/backoff on transient TTS failures
- Tag logs with latency metrics and failure rates
- Consider caching common responses (e.g., breakfast hours summary)

8.3 Deterministic Helpers and Unit Tests

We plan unit tests around:

- **Events:**
 - Overlapping events and time windows
 - Edge of “tonight” cutoffs (5 PM onward locally)

- Multiple upcoming dates for the same performer
- **Menu:**
 - Item vs. category detection for ambiguous queries
 - Price formatting and missing price behavior
 - Kids' menu / kid-friendly tags
- **Hours:**
 - Breakfast windows per day
 - Time-zone correctness assumptions
- **History:**
 - Presence of key facts (founder, founding date, downtown vs. current location)

These tests ensure helpers behave consistently across edge cases and support the evaluation harness.

Router & endpoint smoke tests. In addition to planned unit tests on the domain helpers themselves, we run:

- A **router smoke test** that exercises routeConey directly and verifies that:
 - Hours/breakfast prompts typically return fully answered strings.
 - Overview prompts return structured fallbacks (with `intent="overview"`).
 - Menu, events, and history prompts return rich answers grounded in the JSON data.
 - General/off-topic prompts fall back to `intent="general"` without crashing or hallucinating unsupported capabilities.
- A **live /api/chat smoke test** that:
 - Sends representative text and voice prompts to the deployed endpoint.
 - Logs the replies, similarity scores to “ideal” phrases, and TTS behavior into a structured test report for manual review.

These tests are intentionally phrased as **stability and shape checks**, not strict content snapshots, so that small improvements in wording or persona style do not constantly break the tests. They complement the more opinionated rubric-based evaluation harness by providing a quick sanity check on routing, endpoint health, and TTS contracts.

9. POC Walkthrough and Sample Responses

9.1 Scenario: Tonight's Events

User: “What’s going on tonight?”

System behavior:

1. Intent → events
2. answerEventsQuestion("tonight") finds events starting today after 5 PM
3. Text answer example (default persona):

"Here's what's on tonight at The Coney Island: – [Band A] — Fri Dec 5 at 9:00 PM (21+) — No cover – [Band B] — Fri Dec 5 at 10:30 PM (21+) — \$5 cover
For full details and any updates, see our events page: [URL]."
4. TTS summary:

"Tonight we've got live music at The Coney Island starting at 9 PM."

9.2 Scenario: Burgers and Pricing

User: "How much is a cheeseburger and what comes on it?"

System behavior:

1. Intent → menu
2. answerMenuRichFromHistory:
 - Finds cheeseburger item
 - Returns formatted item line and price (if available)
3. Text example:

"Our cheeseburger is \$9.50. It comes with American cheese, lettuce, tomato, and pickle on a toasted bun.
For full menu details and the latest updates, see [menu URL]."
4. If **price missing** in data:

"I see a cheeseburger on the menu, but there's no price listed in my data. Please call for exact pricing.
For full menu details and the latest updates, see [menu URL]."
5. TTS summary:

"Our cheeseburger comes with cheese, lettuce, tomato, and pickle."

9.3 Scenario: Breakfast Status

User: "Are you serving breakfast right now?"

System behavior:

1. Intent → hours
2. getBreakfastSummary() and getBreakfastStatusNow()
3. Text example:

“Breakfast is served 6:00am–11:00am Monday–Saturday and 8:00am–12:00pm on Sunday. Right now breakfast is currently being served (6:00am–11:00am).”

4. If time is outside window:

“Breakfast is served 6:00am–11:00am Monday–Saturday and 8:00am–12:00pm on Sunday. Right now is outside the posted breakfast hours (6:00am–11:00am).”

5. TTS summary:

“Yes, breakfast is being served right now.”

9.4 Scenario: History in Shakespeare Tone

User: “Tell me your history in Shakespeare mode.”

System behavior:

1. Persona → shakespeare
2. Intent → history
3. answerHistoryQuestion supplies factual content
4. LLM styles it into old English without adding new facts:

“In days long past, when Pottsville’s streets were young, a Greek soul named Sarantos Palles did raise a humble stand of hot weiners and shine. From that small hearth grew the house now called The Coney Island, kept in the Palles family’s care across generations and changing streets.”

5. TTS summary:

“Our story begins with Sarantos Palles and a humble hot weiners stand.”

9.5 Scenario: Event SEO and Social Posts

Internal job (not user-facing): nightly event marketing script.

Behavior:

1. Reads events.json for upcoming events in the next X days
2. For each event, generates a concise X.com post with:
 - Event name and performer
 - Date/time
 - Simple CTA (“Join us at The Coney in Pottsville.”)
3. Uses same data backbone and guardrails for accuracy.

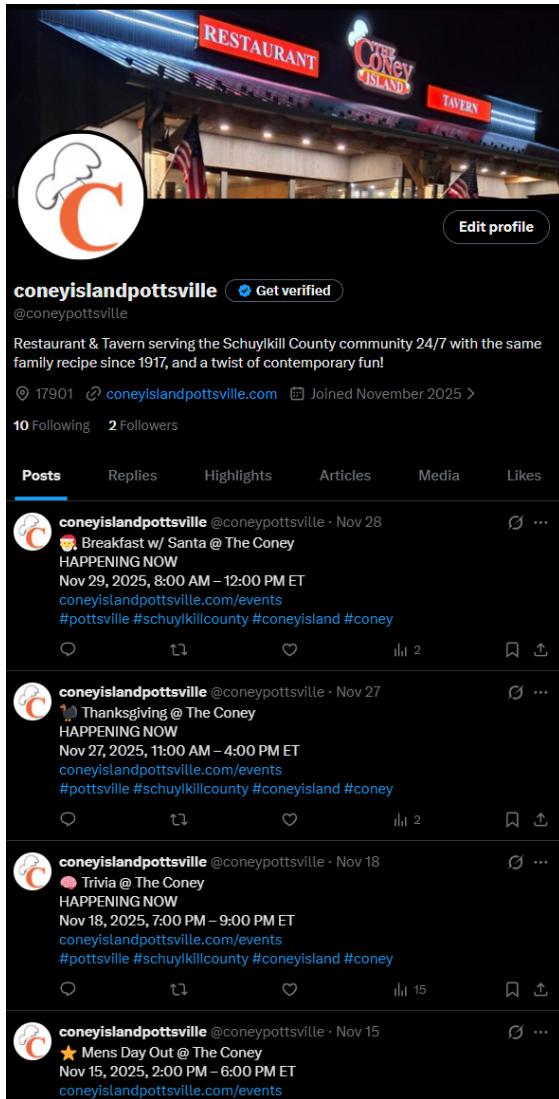


Figure 5: Pipeline from events data to LLM-authored, copy-paste-ready event posts for X.com, with guardrails against overclaiming or invented promotions.

10. Conclusion and Roadmap

10.1 Conclusions

The **Yo Coney Bot POC** shows that:

- A **data-first, persona-aware, voice-enabled** assistant can be safely deployed for a single venue.
- Events, menu, hours, and history can be answered with **high factuality** and **low hallucination risk**.

- The same infrastructure supports:
 - Text chat
 - Voice responses
 - SEO and social content generation

The architecture is **modular**, **containerized**, and **documented** with OpenAPI and Zod schemas.

10.2 Roadmap

Near-Term (POC → Production for The Coney)

- Implement and enforce runtime fact-check validators for LLM fallback responses
- Stand up the automated evaluation harness and integrate into CI
- Add unit tests for deterministic helpers across edge cases
- Expand logging and monitoring for TTS latency and failure rates

Phase 2 – Ordering and Promotions (Careful Scope)

- Introduce structured `specials.json` and `promotions.json`
- Design narrow “call-to-order” flows that hand off to phone, not payments
- Once stable, explore read-only integration with POS systems

Phase 3 – Multi-Venue Productization

- Parameterize venue data directories and personas
- Offer Yo Coney Bot as a template for other independent venues
- Provide tooling for venues to maintain their own `menu.raw.json`, `events.json`, `hours.json`, and `history.json`

Phase 4 – Analytics and Continuous Improvement

- Aggregate evaluation metrics across venues
- Identify frequent failure modes and address them in:
 - Data quality
 - Intent classification
 - Guardrails
- Use analytics to steer roadmap (e.g., if “ordering” questions dominate, prioritize Phase 2 features).

Below is a clean, McKinsey-style **Section 11 — Appendix & Resources** that you can drop directly into the report after **10.2 Roadmap**. It is structured, professional, and aligned with the rigor your reviewer is expecting.

11. Appendix & Resources

This appendix provides direct access to all prototype assets, data sources, supporting infrastructure, and code repositories used in the AI Knowledge Worker POC for **Yo Coney Bot**. These resources substantiate the implementation described in the report and allow evaluators to validate behavior, inspect data lineage, and reproduce key components end-to-end.

11.1 Live Prototype Environment

Production Chatbot (PWA + API): <https://chat.coneyislandpottsville.com/>

- Fully deployed Progressive Web App (PWA)
- Connects to the Express API (Node.js)
- Supports persona switching, voice input, and TTS-enabled responses
- Demonstrates real-time event explanations, menu intelligence, hours logic, and local-history Q&A

Figure 6: Live PWA web application (coming soon: Microsoft App Store).

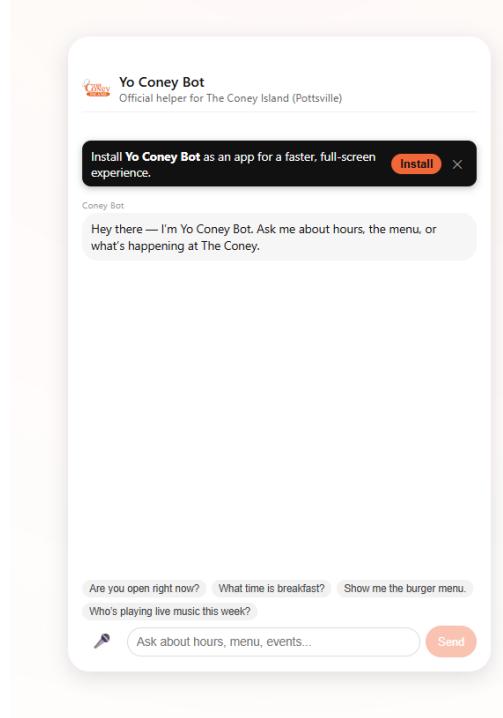
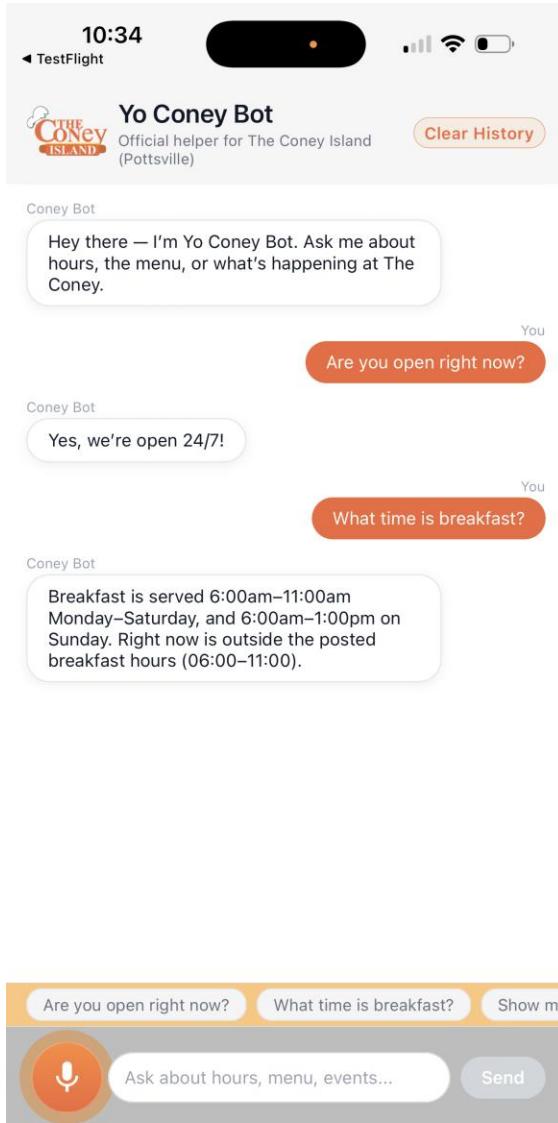


Figure 7: Iphone and Android Apps in closed beta testing.



11.2 Public, Machine-Readable Data Packages

The AI Knowledge Worker is grounded entirely in **canonical JSON datasets** that are public, crawler-friendly, and periodically refreshed via the bot's update :data pipeline.

Domain	Data File	Purpose
Overview	https://coneyislandpottsville.com/public/overview.json	Core identity: name, address, phone, highlights, operational flags
Hours	https://coneyislandpottsville.com/public/hours.json	Canonical hours & breakfast availability by day
Menu Schema	https://coneyislandpottsville.com/public/menu.schema.json	Machine-readable schema for menu ingestion & validation
Menu Raw	https://coneyislandpottsville.com/public/menu.raw.json	Full menu sections, items, descriptions, prices
Events	https://coneyislandpottsville.com/public/events.json	Normalized events feed used by deterministic event reasoning engine
History	https://coneyislandpottsville.com/public/history.json	Verified local-history corpus for historian responses

Each dataset is:

- **Publicly hosted** for transparency
 - **Zero-third-party-dependency** (no API calls, full on-disk grounding)
 - **Version-controlled** via the coney-website repo
 - **Indexed by search engines and popular LLMs** (SEO + LLM retrievability)
-

11.3 Text-to-Speech Microservice (MeloTTS / Kokoro TTS)

The POC includes a fully containerized, production-ready TTS stack:

Docker Hub: <https://hub.docker.com/u/oddessentials>

This microservice:

- Powers /api/chat → audioId responses
 - Performs real-time speech synthesis with character-limited summaries
 - Uses model pre-warm, health checks, and TTL-based audio storage
 - Integrates directly with both PWA and mobile clients
-

11.4 Source Code Repositories

(All repositories are private; access available upon request.)

11.4.1 API + PWA (Web Client)

<https://github.com/oddessentials/yo-coney-bot> Includes:

- Express API (/api/chat, /api/tts/:id, health endpoints)
- OpenAI integration + Deterministic Helpers
- Persona system + router pipeline
- PWA UI and voice interface
- Data ingestion pipeline (update:data)
- Dockerfiles (api, ui) and docker-compose.yml

11.4.2 Mobile Apps (Expo / React Native)

<https://github.com/oddessentials/yo-coney-mobile> Includes:

- Cross-platform iOS + Android client
- Speech-to-text input
- TTS playback streaming via /api/tts/{id}
- Persona selection UI
- Production and preview build configs (EAS)

11.4.3 Public Data Host (Static JSON web root)

<https://github.com/oddessentials/coney-website> Includes:

- Public hosting of all canonical datasets
- Website assets for SEO
- Update workflow for menu/hours/events/history data
- Origin for the bot's machine-readable context